
Transfer Grokking: Linear Projection of Fourier Representations Across Model Residual Streams

Saparmyradov Saparmyrat
RTU MIREA
saparmyrat.saparmyradov@gmail.com

Abstract

Grokking transformers learn perfectly generalizing modular arithmetic circuits that encode discrete Fourier features in their residual stream. Transferring these compiled representations into large language models (LLMs) would enable neural reuse of structured knowledge, but prior attempts have found two independent barriers: (1) mean-squared-error (MSE) projection training produces weight matrices that are perfectly probeable yet misaligned with the target LLM’s unembedding directions, and (2) injecting the projected representation at intermediate layers corrupts the LLM’s remaining computation. We show that both barriers are resolvable. Training the projection via cross-entropy through a frozen unembedding layer achieves perfect logit-lens accuracy (1.0), while injecting at the final layer (L31 for Phi-2) achieves perfect task accuracy (1.0) at full strength and 0.705 at a practical operating point with only mild perplexity degradation (65.9 vs. 63.3 baseline). Results replicate across modular addition and multiplication (both mod 97), demonstrating that the transferred representation is task-agnostic. Our findings suggest that grokked and language-model representations are fundamentally commensurable - the bottleneck is not geometric but operational, and resolvable by targeting the last layer.

1 Introduction

Mechanistically interpretable circuits discovered in small transformers trained on algorithmic tasks—such as the discrete Fourier features that implement modular arithmetic [6, 2]—offer a unique window into how neural networks implement structured computation. These circuits are sparse, precise, and perfectly generalizing. A natural question is whether they can be *transferred* into large language models (LLMs) as reusable computational primitives.

Prior work on transferring representations between neural networks has largely focused on measuring similarity (SVCCA [7], CKA [3]), stitching representations [1, 5], or steering model behavior via activation vectors (CAA [8], activation addition [9], inference-time intervention [4]). These approaches either diagnose alignment or modify LLM behavior, but they stop short of transferring the *entire computational content* of a grokked circuit into an LLM’s residual stream.

We identify and resolve two independent barriers to such transfer:

1. **Loss misalignment.** Training a linear projection $W : \mathbb{R}^{128} \rightarrow \mathbb{R}^{2560}$ between the grokked model’s hidden state and the LLM’s residual stream using MSE produces a weight matrix whose output is probeable (1.0 accuracy) yet decodes to near-random logits (0.012). Training W via cross-entropy through the frozen unembedding layer resolves this, achieving perfect logit-lens accuracy (1.0).

2. **Injection depth.** Even with a correctly trained W , injecting the projected state at intermediate layers (e.g., L10 of Phi-2) degrades accuracy because the remaining computation is corrupted. Injecting at the final layer (L31) avoids this, achieving perfect task accuracy (1.0) with only mild perplexity degradation (65.9 vs. 63.3 baseline at $\alpha = 0.5$).

Both results replicate across modular addition and multiplication (mod 97), confirming the findings are operation-agnostic.

2 Setup

Grokked model. We train a 2-layer transformer ($d_{\text{model}} = 128$, $d_{\text{MLP}} = 512$) on $(a + b) \bmod 97$ with 30% of the 97^2 pairs, using weight decay 1.0 and full precision. Training runs for up to 100,000 epochs with rolling-window grokking detection (10-epoch window of validation accuracy ≥ 0.99). The model uses direct token IDs (0–96) as both input and output, with a vocabulary size of 97. The same architecture is also trained on $(a \times b) \bmod 97$ and groks successfully.

LLM target. We use Microsoft Phi-2¹ ($d_{\text{model}} = 2560$, 32 layers) as the primary target LLM. Phi-2 encodes all numbers 0–96 as single BPE tokens (token IDs 15–6659), making downstream evaluation of modular arithmetic well-defined—unlike other LLMs whose tokenizers split multi-digit numbers into subword tokens, preventing lm-head-based evaluation.

Projection layer. A linear layer $W \in \mathbb{R}^{2560 \times 128}$ (no bias) maps the grokked model’s penultimate-layer hidden states to Phi-2’s residual stream dimension. The prompt template is "# (a op b) % 97 =" for both training and evaluation.

Evaluation metrics. We report four metrics:

1. **Logit lens accuracy:** accuracy of $W(h)$ projected through the frozen lm_head (sliced to the 97 number-token logits).
2. **Probe accuracy:** 97-way LogisticRegression accuracy on $W(h)$.
3. **Patch accuracy:** accuracy after injecting the α -interpolated state $h' = (1 - \alpha) \cdot h_{\text{LLM}} + \alpha \cdot W(h_{\text{grokked}})$ at a given layer.
4. **Perplexity:** WikiText-2 last-token perplexity after patching.

Results are reported as mean $\pm$ std over five random seeds (42–46).

3 Two Barriers to Transfer

The geometry – linear separability gap. Figure 1 (adapted from prior experiments) summarizes the disconnect between representational similarity (cosine similarity) and linear separability (probe accuracy) across several transfer configurations. Small-to-small transfer achieves both perfect similarity and perfect separability (1.0, 1.0). Phi-2-to-Phi-2 shows a gap (1.0, 0.71): activations from a single prompt template are perfectly self-similar ($\cos = 1.0$, trivially) yet only partially linearly separable for the arithmetic task (probe 0.71), confirming that geometric identity does not guarantee full task-relevant structure even within a single model. The Clean Experiment (Small-to-Big, same tokenizer) retains high separability (0.94) despite low cosine similarity (0.30), while the Embed Patch bypasses the tokenizer altogether and achieves high similarity (0.82) but near-zero separability (0.01).

Barrier 1: MSE misalignment. Training W by minimizing the MSE between $W(h_{\text{grokked}})$ and $L10(h_{\text{Phi-2}})$ produces a projection whose output is perfectly probeable by a 97-way logistic regression (1.0) but decodes to near-random logits through the frozen lm_head (0.012). The probe’s success confirms that the arithmetic information is present in $W(h)$ —the bottleneck is coordinate alignment with the unembedding directions, not information loss. MSE treats all 2560 dimensions equally, but the lm_head’s weight vectors define a privileged set of decoding directions. Minimizing pixel-level distance does not align with those directions.

¹<https://huggingface.co/microsoft/phi-2>




Figure 1: **Cosine similarity vs. probe accuracy.** Four transfer configurations reveal a consistent gap between representational similarity and linear separability. The $y = x$ diagonal marks perfect alignment.

Barrier 2: Injection depth. Even with a perfectly aligned W , injecting at intermediate layers (e.g., L10) yields accuracy consistently below baseline for both MSE and CE projections (Table 2 in Appendix A). The injected state overwrites the LLM’s ongoing computation—at L10, 21 layers of remaining processing are corrupted by a signal unrelated to the natural language context.

4 CE Training Resolves Barrier 1

Training procedure. We train $W : \mathbb{R}^{128} \rightarrow \mathbb{R}^{2560}$ by minimizing cross-entropy through the frozen `lm_head`. The `lm_head`’s weight matrix $\text{LM} \in \mathbb{R}^{51200 \times 2560}$ is sliced row-wise to the 97 number-token rows ($\text{LM}_{\text{num}} \in \mathbb{R}^{97 \times 2560}$). The forward pass is:

$$\mathcal{L} = \text{CE}(\text{LM}_{\text{num}} \cdot W(h_{\text{grokked}}), y)$$

where y is the ground-truth answer (0–96). W is trained with AdamW ($\text{lr} = 10^{-3}$, $\text{wd} = 10^{-2}$) for 5,000 epochs on the same 70/30 train-test split. No layer-normalization layers are included—the gradient flows directly through the linear chain $W \rightarrow \text{LM}$, ensuring W must learn to produce states along the `lm_head`’s decoding directions.




Figure 2: **MSE vs. CE projection.** CE training achieves perfect logit-lens accuracy (1.0) while MSE training yields near-random decoding (0.012). Both projections are perfectly probeable (1.0). The dashed line marks the Phi-2 baseline (0.235).

Results. CE-trained W achieves perfect logit-lens accuracy (1.0) on the held-out test set, resolving Barrier 1 completely. The probe on $W(h)$ also achieves 1.0, confirming no information is lost.

Crucially, W_{CE} learns to produce states along the `lm_head`’s decoding directions, not along the Phi-2 residual’s natural basis—cosine similarity with L10 target activations is effectively zero (-0.0005). This confirms that the MSE objective was the sole cause of the logit-lens misalignment; the information was never missing, only misdirected.

For multiplication mod 97, CE training likewise achieves W_{CE} with logit-lens accuracy 1.0, demonstrating the approach is operation-agnostic.

5 L31 Injection Resolves Barrier 2

Motivation. If the remaining layers after injection corrupt the signal, then injecting at the *last layer*—where no further computation remains—should avoid this corruption entirely. We test this hypothesis by injecting $h' = (1 - \alpha) \cdot h_{\text{LLM}}^{(L)} + \alpha \cdot W_{\text{CE}}(h_{\text{grokked}})$ at the final transformer block (L31 for Phi-2), sweeping $\alpha \in \{0.0, 0.3, 0.5, 0.7, 1.0\}$.




Figure 3: **L31 α -sweep for addition mod 97.** Blue line: accuracy (left axis). Red dashed line: perplexity (right axis, log scale). Gray dashed line: random baseline ($1/97 \approx 0.010$). Shaded regions: ± 1 std over 5 seeds (negligible at reported precision for accuracy). Accuracy rises monotonically with α , while perplexity degrades only mildly at $\alpha = 0.5$.

Table 1: L31 patch accuracy (α -sweep, mean $\pm$ std over 5 seeds).

α	Addition	Multiplication
0.0	0.235 ± 0.000	0.010 ± 0.000
0.3	0.490 ± 0.000	0.080 ± 0.000
0.5	0.705 ± 0.000	0.370 ± 0.000
0.7	0.995 ± 0.000	0.930 ± 0.000
1.0	1.000 ± 0.000	1.000 ± 0.000

Addition results. At $\alpha = 1.0$, the injected state completely replaces the L31 residual and accuracy reaches 1.0 (Table 1). Performance rises monotonically with α —there is no intermediate degradation. The baseline ($\alpha = 0.0$) is Phi-2’s unmodified accuracy on the mod 97 prompt (0.235), which arises from an implicit bias toward single-digit answers.

Critically, the perplexity tradeoff is manageable. At $\alpha = 0.5$ (the operating point that balances arithmetic accuracy and language modeling quality), perplexity on WikiText-2 is 65.9 ± 0.01 , only

+2.6 above the no-patch baseline of 63.3. At $\alpha = 1.0$, perplexity degrades catastrophically ($\sim 6 \times 10^{14}$), confirming that the injected state is unrelated to the natural language context. In practice, a moderate α (e.g., 0.5–0.7) achieves high arithmetic accuracy (0.71–0.99) while preserving text quality.

Multiplication results. All findings replicate for multiplication mod 97 (Figure 4, Appendix). W_{CE} achieves 1.0 logit-lens accuracy, and L31 injection at $\alpha = 1.0$ yields perfect accuracy (1.0). The intermediate- α accuracy is lower than for addition (0.37 vs. 0.71 at $\alpha = 0.5$), consistent with multiplication’s more complex Fourier structure requiring stronger injection to dominate the residual. Perplexity degrades similarly (66.9 at $\alpha = 0.5$ vs. 63.3 baseline).

Comparison with L10 injection. For contrast, we also evaluate the L31-patched model against the original L10-patch approach from prior work. At L10, CE-trained W actually *underperforms* MSE-trained W (0.26 vs. 0.305 at $\alpha = 0.5$), confirming that the injection depth—not the projection quality—was the primary confound in prior failures. At L31, the remaining computation is zero, so any well-aligned injection signal dominates the residual.

6 Multi-task Validation

The experiments above establish that two barriers—MSE misalignment and injection depth—independently prevent transfer, and that both are resolvable. To rule out task-specific confounds, we replicate the entire pipeline on $(a \times b) \bmod 97$:

- **Grokking:** the small model groks multiplication successfully, confirming the training protocol extends to more complex operations.
- **CE projection:** W_{CE} achieves logit-lens accuracy 1.0 (seed 42: 0.025 training loss, 100% validation accuracy at epoch 5,000).
- **L31 injection:** $\alpha = 1.0$ achieves 1.0 accuracy; $\alpha = 0.5$ achieves 0.37 accuracy with perplexity 66.9.

The qualitative pattern is identical across tasks: CE resolves Barrier 1, L31 resolves Barrier 2. The lower intermediate- α accuracy for multiplication is quantitatively consistent with the operation’s higher Fourier complexity and does not indicate a structural difference.

7 Discussion

Two independent barriers. Our results decompose the general problem of “representation transfer between models” into two orthogonal subproblems: (1) aligning the projection with the target model’s decoding directions, and (2) choosing the correct injection depth. Each barrier admits a distinct solution, and neither alone is sufficient.

Why CE works and MSE fails. MSE is a dimension-uniform loss—it treats all 2560 features of the target residual stream equally. The `lm_head`, however, defines a small set of privileged directions (97 out of 2560 for our task) that determine decoding. A projection trained with MSE learns to reconstruct the full distribution of activations, not to align with these privileged directions. CE through the frozen `lm_head` forces W to discover exactly those directions, because any deviation from them directly increases classification loss. This explains the stark gap in logit-lens accuracy (0.012 vs. 1.0) despite identical probe performance.

Why the last layer. Injecting at L10 degrades the LLM’s remaining computation because the injected state carries task-specific Fourier structure unrelated to the natural language context. Every subsequent layer attempts to integrate this foreign signal into its language-modeling computation, resulting in corrupted representations. At the last layer (L31), no further computation remains—the residual stream is projected directly to logits. An injection here replaces the final state without cascading effects. This mirrors the “neural function call” pattern: the last layer acts as a dispatch point where computation from another model can be spliced in.

Generality and limitations. Our findings hold across two modular arithmetic operations (addition and multiplication) and across five random seeds. The Phi-2 tokenizer’s property of encoding

all 0–96 as single tokens was essential for lm-head-based evaluation—models with number-splitting BPE (e.g., Qwen2-Math-1.5B, which splits 87/97 numbers) require alternative evaluation strategies (Appendix F).

The primary limitation is the moderate accuracy at intermediate α —mixing strengths that preserve language quality (e.g., $\alpha = 0.5$) only achieve 0.71 accuracy for addition and 0.37 for multiplication. This is a genuine tradeoff: the injected Fourier features are unrelated to the natural language context, so strong injection necessarily degrades general language modeling. Applications requiring near-perfect arithmetic accuracy would use $\alpha \geq 0.7$, accepting a higher perplexity cost.

Broader implications. Our results demonstrate that grokked and language-model representations are fundamentally commensurable—the bottleneck is operational, not geometric. This suggests a general principle: transferring a representation between models requires (1) aligning the projection with the target’s output decoder, not its hidden-state geometry, and (2) injecting at a depth where no further target-model computation remains. We conjecture that this principle generalizes beyond modular arithmetic to other structured tasks (e.g., board-game evaluation, symbolic reasoning) and other model pairs.

Acknowledgments

We thank the NeurIPS community for feedback and the developers of the transformer-circuits library and HuggingFace Transformers for making this work possible.

References

- [1] Yamini Bansal, Preetum Nakkiran, and Boaz Barak. Revisiting model stitching to compare neural representations. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- [2] Bilal Chughtai, Lawrence Chan, Neel Nanda, Jessica Z. Forde, and David Ha. A toy model of universality: Reverse engineering how networks learn group operations. In *International Conference on Machine Learning (ICML)*, 2023.
- [3] Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey E. Hinton. Similarity of neural network representations revisited. In *International Conference on Machine Learning (ICML)*, 2019.
- [4] Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [5] Jack Merullo, Carsten Eickhoff, and Ellie Pavlick. Circuit component reuse across tasks in transformer language models. In *International Conference on Learning Representations (ICLR)*, 2024.
- [6] Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. Progress measures for grokking via mechanistic interpretability. In *International Conference on Learning Representations (ICLR)*, 2023.
- [7] Maithra Raghu, Justin Gilmer, Jason Yosinski, and Jascha Sohl-Dickstein. SVCCA: Singular vector canonical correlation analysis for deep learning dynamics and interpretability. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [8] Nina Rimskey, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Matt Turner. Steering llama 2 via contrastive activation addition. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- [9] Alexander Matt Turner, Lisa Thiergart, David Udell, Gavin Leech, Ulisse Mini, and Monte MacDiarmid. Activation addition: Steering language models without optimization. In *CoRR*, volume abs/2308.10248, 2023.

Appendices

A L10 Alpha Sweep

For completeness, Table 2 shows the accuracy of MSE- and CE-trained projections when injected at L10 of Phi-2, confirming that neither achieves meaningful transfer at this depth.

Table 2: L10 patch accuracy (α -sweep, seed 42).

α	MSE (add)	CE (add)	CE (mult)
0.0	0.235	0.235	0.010
0.3	0.290	0.265	0.010
0.5	0.305	0.260	0.010
0.7	0.280	0.230	0.005
1.0	0.015	0.040	0.020

B Multiplication Alpha-Sweep Figure




Figure 4: **L31 α -sweep for multiplication mod 97.** Same format as Figure 3. Accuracy reaches 1.0 at $\alpha = 1.0$ with a steeper climb after $\alpha = 0.5$ compared to addition.

C Cross-task Comparison

Figure 5 compares addition and multiplication accuracy side-by-side, highlighting the consistent pattern.

D Failure-mode Details

SVCCA + Noise Injection. Prior work used SVCCA with $k = 20$ to compare residual-stream representations. Raw CCA on 128-dimensional (Small) vs. 512-dimensional (Big) activations over $N = 2823$ test points overfits to ~ 1.0 ; dimensionality reduction to $k = 20$ gives meaningful alignment scores. Layer-by-layer SVCCA reveals that layers align by position, not cross-functionally (e.g., Small[1] $\leftrightarrow$ Big[5] = 0.835).

Noise injection experiments calibrated to the embedding norm (~ 22.65) used $\sigma \in \{0.05, 0.10, 0.20, 0.50\}$ —higher σ values (0.5, 1.0, 2.0) completely destroy signal. Steering vec-




Figure 5: **Cross-task comparison.** Addition (blue) and multiplication (orange) accuracy across α values. Both tasks reach 1.0 at $\alpha = 1.0$; multiplication lags at intermediate α values.

tors become indistinguishable from random directional vectors when cosine similarity drops below ~ 0.7 .

Clean Experiment (Small $\rightarrow$ Big). A linear projection was trained from Small ($d = 128$) to Big ($d = 512$) using MSE with the same tokenizer (direct token IDs 0–96). Despite matching tokenizers, testing on natural-language prompts failed, establishing that tokenizer mismatch is not the primary barrier to transfer.

Embed Patch. Bypassing Phi-2’s BPE tokenizer by inserting $W_{\text{emb}}(h_{\text{grokked}})$ directly as inputs_embeds achieves cosine similarity 0.82 with the natural embeddings but near-zero task accuracy (0.01). This confirms that representational similarity does not imply functional transfer.

Multi-layer injection. Injecting the projected state at 5 layers simultaneously (L2, L8, L14, L20, L26) with per-layer projection matrices W_{layer} degrades accuracy below the single-layer L10 baseline (0.105 at $\alpha = 0.3$). Per-layer W matrices are approximately equal to a single shared W , confirming that layer-specific alignment does not matter.

Nonlinear adapter. A two-layer MLP ($2560 \rightarrow 2560 \rightarrow 2560$) between $W(h)$ and the frozen lm_head achieves the same accuracy as a linear adapter (1.0). Since two linear layers compose into one linear transformation, this confirms the bottleneck is not in the projection’s expressivity.

E Single-template Probe on L31

A logistic regression probe (Linear $2560 \rightarrow 97$) trained on Phi-2’s L31 activations from a single prompt template (“# (a + b) % 97 =”) achieves 0.71 accuracy—the same value as the Phi-2-to-Phi-2 probe in Figure 1—well above random (0.01) but below the grokked model’s 1.0. This confirms that L31 activations are partially linearly separable under the standard template, though template sensitivity remains critical: generalization to diverse natural language templates is substantially weaker (Appendix D).

F Cross-model Probe (Bypassing Tokenizer)

For models with number-splitting BPE (e.g., Qwen2-Math-1.5B, which splits 87/97 numbers), lm_head -based evaluation is impossible because only 10 unique token outputs exist for 97 classes. We bypass this by training a logistic regression probe on W_{CE} -injected last-layer hidden states. At $\alpha = 1.0$, both Phi-2 L31 and Qwen2-Math-1.5B L27 achieve probe accuracy ≥ 0.999 , confirming that the injected Fourier features dominate the residual regardless of target model architecture.

G Per-seed Data

All five seeds (42–46) produce indistinguishable accuracy values (zero variance at reported precision) for both addition and multiplication. Perplexity variance across seeds is negligible ($\sigma < 0.01$).

PPL at $\alpha = 0.5$). Raw weight matrices differ numerically (max element-wise difference ~ 0.7 for addition, ~ 1.3 for multiplication), confirming that the zero accuracy variance is a property of the 97-class decision boundary, not a seed propagation error.

Table 3: Per-seed perplexity at $\alpha = 0.5$.

Seed	Addition PPL	Multiplication PPL
42	65.88	66.88
43	65.89	66.88
44	65.88	66.88
45	65.87	66.88
46	65.91	66.88
Mean $\pm$ Std	65.89 $\pm$ 0.01	66.88 $\pm$ 0.00